

Manufactured Gradients: State AI Exposure and US Business AI Adoption at DeepSeek-R1

Hauke Hillebrandt*

July 2026

analysis pass of record: 2026-07, `natex` v0.2.0, seed 0

Abstract

US states more exposed to AI — measured by the Anthropic Economic Index (AEI) usage-per-capita index — did not detectably accelerate business AI adoption relative to less-exposed states at the release of DeepSeek-R1 on 20 January 2025. In Census Bureau Business Trends and Outlook Survey (BTOS) state panels (2,727 state-wave cells, 51 states, 71 biweekly waves, 2023–2026), the change in each state’s adoption-trend slope at R1 rises by +1.635 pp/yr per AEI index unit (HC1 se 1.458) — reproducing the scout estimate to the third decimal — but is statistically null: exposure-permutation $p = 0.257$. The same statistic at non-event placebo cutoffs brackets the headline (+1.56 at $t = 2024.30$, +2.24 at 2024.50), and a placebo *outcome* — remote work, whose federal return-to-office executive order was signed the same day as R1 — matches it (+1.97), with the weighted placebo version itself nominally significant ($p = 0.045$): the design manufactures gradients. Every nominally significant estimate in the battery traces to a known mechanism. Long-window negative gradients (−4.31 at $w = 1.0$) vanish under old-wording-only data — a splice artifact, because the nationally calibrated 1.553 rewording ratio misprices high-adoption, high-AEI states (corr +0.676). The two-way fixed-effects level effect (+0.527 pp, permutation $p = 0.011$) is a pre-trend: the high-minus-low gap was already growing at +0.600 pp/yr before R1, and state-specific trends flip the estimate to −0.38. The scan’s localization at the first post-R1 wave is a mechanical rediscovery of the constructed treatment column. There is no credible evidence that high-AEI-exposure states accelerated AI adoption differentially at DeepSeek-R1: the BTOS state Bartik-style design, excluded from an earlier screen for want of exposure weights, stays dead even with real ones.

1 Introduction

The release of the DeepSeek-R1 reasoning model on 20 January 2025 [7] is the natural candidate event for a break in US business AI adoption: a companion sector-level analysis [10] of the Census Bureau’s Business Trends and Outlook Survey (BTOS) found a sharp acceleration in AI-exposed sectors at exactly the first post-R1 survey wave, but attribution to R1 failed a placebo battery. This paper asks the cross-sectional version of the same question: did *states* whose workforces use AI more intensively accelerate measured business AI adoption differentially at R1? The design — state-level BTOS adoption interacted with an external AI-exposure measure — was excluded from an earlier automated screen for want of exposure weights; here we revive it with real ones.

*University College London. Email: `ucjthhi@ucl.ac.uk`. This paper and the underlying `natex` software were prepared with substantial assistance from Anthropic’s Claude models; the author reviewed the analyses and text and is responsible for all remaining errors.

Both sides of the design rest on new measurement. The outcome is the BTOS AI use rate, a biweekly probability-sample estimate of the share of US firms using AI, introduced and validated by [3]. The exposure is the Anthropic Economic Index (AEI), which maps Claude conversations to economic tasks [8]; its geographic report [1] publishes a usage-per-capita index by US state, which we take as the cross-state exposure measure. Interacting differential exposure with a common shock is the logic that shift-share (Bartik) designs formalize [6], and its known failure modes discipline the analysis: exposure here is measured *after* the outcome window, so every estimate below is a descriptive gradient, never a causal Bartik coefficient; pre-existing differential trends can masquerade as treatment effects [13]; and slope changes at a single date in smooth aggregate series are spuriously significant often enough that placebo distributions at non-event dates, in the spirit of [5], are the appropriate yardstick. The per-state slope-change contrast is a kink-style estimand in the sense of the difference-in-kinks design [2], estimated and falsified with the `natex` toolkit [12], whose scan-based DiD [9] supplies a localization check. The battery converts one reproduced headline gradient and three nominally significant side results into a clean null plus three identified artifacts.

2 Data

The outcome panel combines three BTOS sources. The state sheets of the AI core workbook (DRB approval CBDRB-FY25-0425) give biweekly estimates and standard errors of the share answering “Yes” to Question 7 (“In the last two weeks, did this business use Artificial Intelligence (AI) in producing goods or services?”), waves 202319–202520. A fresh census.gov download extends the panel under the December-2025 rewording (“...in any of its business activities?”), waves 202524–202614; new-wording values are divided by 1.553, the national new/old measurement-RDD ratio estimated at the rewording cutoff by a companion paper [11], with an additive -6.04 pp variant as a diagnostic. The archived remote-work Question 6 (share of businesses with paid employees working from home at least one workday), waves 202417–202514, supplies a placebo outcome (1,221 state-wave cells). The assembled panel has 2,727 state-wave cells across 51 states (50 plus DC) and 71 waves, $t = 2023.710$ – 2026.507 in fractional years of collection-period midpoints; 24.7% of the state-by-wave grid is disclosure-suppressed, concentrated in small states.

Exposure is the AEI `usage_per_capita_index` per state [1], averaged over the 2026-04-01..05-01 and 2026-05-01..06-01 data windows (verified: $CA = (1.56 + 1.62)/2 = 1.59$); it ranges from 3.52 (DC) to 0.245 (WV). Two features bind the interpretation. First, the exposure window (April–June 2026) lies entirely *after* the outcome window, so exposure is endogenous by construction and no estimate below is causal, whatever its p -value. Second, the index measures Claude usage specifically, proxying all-AI exposure with attenuation and tech-hub confounding both plausible.

3 Design and Methods

Exposure gradient at a known cutoff. For each state s , local OLS slopes of the adoption share in calendar time are fit on each side of $t_0 = 2025.055$ (DeepSeek-R1; first post wave 202503) within $\pm w$ years, and the slope change $\Delta\kappa_s$ is the post-minus-pre difference — the right-minus-left kink convention of [2], with the cross-state gradient replacing the group contrast. The headline statistic is the cross-state regression of $\Delta\kappa_s$ on AEI exposure, b in pp/yr per index unit. The spec of record is $w = 0.5$, unweighted, restricted to states with at least 6 usable waves per side ($n = 36$; DC never enters — it is suppressed below the wave minimum). Inference is primary by exposure permutation (AEI values reassigned across states, $B = 10,000$, `default_rng(0)`), with

HC1 standard errors reported alongside [4]; windows $w \in \{0.5, 0.75, 1.0, 1.4\}$ are examined, and an SE-weighted variant fits $1/\text{se}^2$ WLS slopes from the published State SE sheets and combines them by inverse-variance meta-analysis.

Falsification battery. (i) placebo cutoffs at the non-event dates 2024.30 and 2024.50, at o1 (2024.6967, a real secondary event), at 2025.30, and at 2025.555 (a window that straddles the rewording splice); (ii) the placebo outcome (remote work) at the R1 cutoff — noting that the federal return-to-office executive order was signed on 2025-01-20, the *same day* as R1, so remote-work breaks at that date are expected for non-AI reasons; (iii) leverage checks dropping DC and CA, and a Spearman rank gradient; (iv) splice diagnostics rerunning the long-window gradients under the multiplicative ($/1.553$) splice, the additive (-6.04 pp) splice, and old-wording-only data ($t < 2025.8$).

Panel estimators. A tercile two-way fixed-effects (TWFE) level DiD interacts the top AEI tercile (17 states, cut 1.002) with post-R1, with wave and state fixed effects, CR1 standard errors by state [4], a placebo-in-space permutation of the treated set ($B = 2,000$), and a state-specific linear-trends variant. The **natex** survey pipeline’s scan-based DiD (SuDDDS) [9] searches break dates and windows on the same panel (seed 0, normal model, AR(1)-by-unit permutation null) as a localization and role-assignment check.

Honest-inference notes, stated as run. HC1/CR1 intervals on smooth biweekly aggregates are potentially oversized, so permutation p is the primary inference throughout. The **natex** refusals are reported, not patched over: the SuDDDS anticipation test refused (Holm p all NaN), and all three effect estimators (dd, synthetic control, GESS) refused with “only 0 usable placebos; ≥ 5 required” on this unbalanced panel — the manual TWFE and permutation estimates below are the effect numbers of record.

4 Results

The scout number reproduces exactly, and it is null. At the spec of record ($w = 0.5$, unweighted, 36 states), the R1 slope-change gradient is $b = +1.635$ pp/yr per AEI index unit (HC1 se 1.458; classical se 1.428, matching the scout’s “se 1.43, t 1.15”; $t_{\text{HC1}} = 1.12$; correlation 0.193), with exposure-permutation $p = 0.257$ (Table 1, panel A; Figure 1a). The SE-weighted version is $+1.877$ (HC1 se 1.922, $p_{\text{perm}} = 0.360$); the Spearman rank gradient is $\rho = 0.098$ ($p = 0.568$); dropping DC and CA gives $+1.54$ ($p = 0.326$) unweighted and $+0.53$ ($p = 0.804$) weighted. DC — the highest-exposure unit at AEI 3.52 — never enters the gradient sample in the first place, so no single unit drives the (non-)result.

Placebos bracket the headline. The identical statistic at non-event cutoffs returns $+1.56$ ($p = 0.277$) at $t = 2024.30$ and $+2.24$ ($p = 0.162$) at $t = 2024.50$ — magnitudes that *bracket* the R1 estimate — and -0.23 ($p = 0.889$) at o1 (Table 1, panel B; Figure 1b). The remote-work placebo outcome at R1 is $+1.97$ ($p = 0.288$), the same magnitude as the headline; its SE-weighted version is $+4.39$ (HC1 se 1.895) with permutation $p = 0.045$ — a nominally significant *placebo*. A design that produces R1-sized gradients at dates where nothing happened, and a significant gradient for an outcome whose own policy shock (the federal return-to-office executive order, signed 2025-01-20) merely shares R1’s date, manufactures gradients; the null at R1 carries no information beyond that.




Figure 1: (a) State-level slope change in BTOS AI use at DeepSeek-R1 ($w = \pm 0.5$ yr, $n = 36$ states with ≥ 6 waves per side) against AEI usage-per-capita exposure, with the unweighted OLS fit: $b = +1.635$ pp/yr per index unit, exposure-permutation $p = 0.257$. (b) The same gradient statistic at non-event placebo cutoffs, at o1, at R1, and for the remote-work placebo outcome at R1, with 95% HC1 intervals; filled markers are nominally significant by permutation at the 5% level. The only filled marker in the battery is a *placebo* — the SE-weighted remote-work gradient at R1 (+4.39, $p = 0.045$).

Artifact 1: the long-window sign flip is the splice. Once the window reaches across the December-2025 rewording, the gradient turns significantly negative: -4.31 ($p = 0.011$) at $w = 1.0$ and -2.53 ($p = 0.011$) at $w = 1.4$ (Table 1, panel C). This is a measurement artifact, identified three ways. AEI exposure correlates +0.676 with the pre-R1 adoption level, so the nationally calibrated multiplicative ratio 1.553 misprices high-level (equivalently high-AEI) states; the additive -6.04 pp splice — wrong in the opposite direction — attenuates the same estimates to -2.92 and -1.07 ; and old-wording-only data collapse both to null (-0.87 , $p = 0.279$ at $w = 1.0$; $+0.44$, $p = 0.513$ at $w = 1.4$). A placebo cutoff at 2025.555, placed so the window straddles the splice with no event nearby, yields -9.16 ($p = 0.00025$) multiplicative and still -5.81 ($p = 0.015$) additive: *neither* national correction is valid cross-state, so state-level gradients through the rewording — including any June-2026 extension of this design — are unidentified.

Artifact 2: the TWFE level effect is a pre-trend. The tercile TWFE DiD gives $\hat{\tau} = +0.527$ pp (CR1 se 0.210) with placebo-in-space permutation $p = 0.011$ ($+0.547$, $p = 0.038$ on the full panel) — nominally the strongest result in the study (Table 1, panel D). But the high-minus-low AEI gap was already growing at $+0.600$ pp/yr before R1 (HC1 se 0.180, 36 pre-R1 waves), which mechanically predicts roughly $+0.3$ pp of the estimate over the post window, and adding state-specific linear trends flips the $w = 0.5$ estimate to -0.38 (se 0.31). This is the textbook pre-trend failure [13]: the level DiD is descriptive only, and the trend-robust estimand — the slope-change gradient — is exactly the null of panel A.

Artifact 3: the scan localization is treatment rediscovery. The SuDDDS scan (seed 0, both panels) maximizes at LLR 154.5 for the empty subset — *all* states — at $t_0 = 2025.0877$,

Table 1: Exposure gradients, placebos, and panel estimates. Gradient rows: slope-change gradient b (pp/yr per AEI index unit) with HC1 se and exposure-permutation p ($B = 10,000$; splice diagnostics $B = 4,000$). TWFE rows: level effect (pp) with CR1 se by state and placebo-in-space permutation p ($B = 2,000$).

Specification	Estimate	se	p_{perm}	n
<i>Panel A: AEI gradient at DeepSeek-R1 ($w = 0.5$ yr)</i>				
unweighted (spec of record)	+1.635	1.458	0.257	36
SE-weighted (WLS + IV meta)	+1.877	1.922	0.360	36
unweighted, drop DC+CA	+1.544	1.651	0.326	35
SE-weighted, drop DC+CA	+0.531	2.316	0.804	35
<i>Panel B: placebo cutoffs and placebo outcome ($w = 0.5$, unweighted)</i>				
placebo cutoff $t_0 = 2024.30$	+1.560	1.224	0.277	33
placebo cutoff $t_0 = 2024.50$	+2.238	1.604	0.162	34
o1 cutoff $t_0 = 2024.6967$	-0.233	1.610	0.889	33
remote work (Q6) at R1	+1.970	1.458	0.288	51
remote work (Q6) at R1, SE-weighted	+4.392	1.895	0.045*	51
<i>Panel C: long windows across the 2025-12 rewording (unweighted)</i>				
$w = 1.0$, spliced (/1.553)	-4.305	1.676	0.011*	39
$w = 1.0$, additive splice (-6.04 pp)	-2.922	1.885	0.020*	39
$w = 1.0$, old wording only	-0.873	1.000	0.279	39
$w = 1.4$, spliced	-2.535	0.945	0.011*	40
$w = 1.4$, old wording only	+0.439	0.645	0.513	40
splice-window placebo $t_0 = 2025.555$, spliced	-9.161	2.150	0.00025*	39
splice-window placebo, additive	-5.806	2.639	0.015*	39
<i>Panel D: tercile TWFE level DiD (pp)</i>				
TWFE, $w = 0.5$	+0.527	0.210	0.011*	964
TWFE, full panel	+0.547	0.236	0.038*	2727
TWFE + state trends, $w = 0.5$	-0.382	0.315	—	964
pre-R1 high-minus-low gap slope (pp/yr)	+0.600	0.180	—	36 waves

Notes: * nominally significant at the 5% level by permutation. The $w = 1.4$ additive-splice gradient is -1.069 ($p = 0.095$). Panel D’s TWFE p -values are placebo-in-space; the state-trends row and the pre-trend slope row report CR1 and HC1 standard errors respectively, with no permutation test run. Every nominally significant row is attributed to an identified artifact in the text: splice mispricing (panel C), a pre-existing differential trend (panel D), or a same-day policy shock to the placebo outcome (panel B).

the first post-R1 wave, with scan $p = 0.010$ under the AR(1)-by-unit null. Because the declared treatment column is the constructed `high_aei×post` indicator, a scan that lights up for all states at the post boundary is a mechanical rediscovery of that column, not outcome evidence — the same role-assignment failure documented in the companion sector study [10]. The composition check passes ($p = 0.707$); the anticipation test and all three effect estimators refused, as stated in Section 3.

5 Caveats and Conclusion

Four caveats bound the claims. First, *exposure is post-outcome*: the AEI index is measured April–June 2026, after every outcome wave used at the R1 cutoff, so even a robust gradient would have been descriptive, not Bartik-causal — the null is about this gradient design, not proof of zero differential adoption. Second, *exposure is Claude-specific*: as a proxy for all-AI exposure it is attenuated and plausibly confounded with tech-hub geography. Third, *coverage*: 24.7% of the state-wave grid is disclosure-suppressed, concentrated in small states; estimation samples run 33–40 states, and DC — the most exposed unit — is never estimable at the primary window. Fourth, *the calendar*: R1’s release week was inauguration week, sharing its date with the federal return-

to-office executive order (2025-01-20) and adjoining the Stargate announcement (2025-01-21), the AI-diffusion export rule (2025-01-13), and a BTOS sample-year rollover, so even a surviving gradient could not have been attributed to R1 alone.

The conclusion is a disciplined null. The headline gradient reproduces exactly and is statistically indistinguishable from the same statistic at dates where nothing happened; the one nominally significant gradient in the $w = 0.5$ battery belongs to a placebo outcome with its own same-day shock; and all three seemingly significant side results — the long-window negative gradients, the TWFE level effect, and the scan localization — dissolve under, respectively, old-wording-only data, state trends, and role-assignment scrutiny. These data provide no credible evidence that high-AEI-exposure states accelerated business AI adoption differentially at DeepSeek-R1, and the state-level splice problem implies that extending the design through the rewording cannot rescue it.

Reproducibility. All estimates were produced with `natex` v0.2.0 [12] at seed 0 from the frozen BTOS workbooks (published at <https://www.census.gov/hfp/btos>) and the public AEI state release [1]; no dataset files are committed. Figure 1 regenerates deterministically from the committed `figures/make_fig.py`, which asserts the headline gradient and the significant-placebo numbers against the numbers of record before drawing.

References

- [1] Appel, R., McCrory, P., Tamkin, A., McCain, M., Neylon, T., and Stern, M. (2025). Anthropic Economic Index report: Uneven geographic and enterprise AI adoption. arXiv:2511.15080.
- [2] Böckerman, P., Jysmä, S., and Kanninen, O. (2025). *Difference-in-Kinks Design*. IZA Discussion Paper No. 18313. <https://docs.iza.org/dp18313.pdf>
- [3] Bonney, K., Breaux, C., Buffington, C., Dinlersoz, E., Foster, L., Goldschlag, N., Haltiwanger, J., Kroff, Z., and Savage, K. (2024). Tracking firm use of AI in real time: A snapshot from the Business Trends and Outlook Survey. NBER Working Paper No. 32319.
- [4] Cameron, A. C., and Miller, D. L. (2015). A practitioner’s guide to cluster-robust inference. *Journal of Human Resources*, 50(2), 317–372.
- [5] Ganong, P., and Jäger, S. (2018). A permutation test for the regression kink design. *Journal of the American Statistical Association*, 113(522), 494–504.
- [6] Goldsmith-Pinkham, P., Sorkin, I., and Swift, H. (2020). Bartik instruments: What, when, why, and how. *American Economic Review*, 110(8), 2586–2624.
- [7] Guo, D., Yang, D., Zhang, H., et al. (2025). DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature*, 645(8081), 633–638.
- [8] Handa, K., Tamkin, A., McCain, M., et al. (2025). Which economic tasks are performed with AI? Evidence from millions of Claude conversations. arXiv:2503.04761.
- [9] Herlands, W., McFowland III, E., Wilson, A. G., and Neill, D. B. (2018). Automated local regression discontinuity design discovery. In *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD ’18)*.
- [10] Hillebrandt, H. (2026). A sharp bend, no verdict: Difference-in-kinks tests of US business AI adoption at DeepSeek-R1. Companion paper, natex paper collection. <https://haukehillebrandt.github.io/natex/btos-sector-did-r1/>
- [11] Hillebrandt, H. (2026). The question is the treatment: A known-cutoff regression discontinuity at the BTOS AI-question rewording. Companion paper, natex paper collection. <https://haukehillebrandt.github.io/natex/btos-rewording-rdd/>
- [12] Hillebrandt, H. (2026). *natex: automated natural-experiment discovery and estimation* (version 0.2.0). Software. <https://github.com/HaukeHillebrandt/natex>
- [13] Roth, J. (2022). Pretest with caution: Event-study estimates after testing for parallel trends. *American Economic Review: Insights*, 4(3), 305–322.